

02_LASSO_example

June 25, 2019

1 Using the LASSO for inference and prediction

- Stepwise linear regression is time consuming - the p-values are unstable and are path-dependent
- Would be nice to have the optimal subset of variables selected in an unbiased way
- Use LASSO/Ridge/ElasticNet for wide data
- Can use LASSO for unbiased feature selection
- LASSO uses regularization which penalizes high coefficients
- Lambda parameter analogy to garden hose - most important features start to trickle through as you open the valve

```
In [1]: # Package "glmnet" contains the LASSO function
        # install.packages( 'glmnet' )
        library( glmnet )

        # Package "mice" is for imputation of missing data
        #install.packages( 'mice' )
        library( mice )

        # Metapackage "tidyverse" imports libraries
        # for data manipulation (dplyr) and plotting (ggplot2)
        library( tidyverse )

        # Package "readxl" has the read_excel() function
        library( readxl )

        # Package "skimr" has excellent descriptive statistics
        # function skim_to_wide()
        library( skimr )

        # Package "GGally" has ggplot2-style scatterplot matrices
        library( GGally )

        # Package "ggfortify" has ggplot2-style regression diagnostic plots
        #install.packages( 'ggfortify' )
        library( ggfortify )
```

```

# Package tictoc has functions to time function calls
library( tictoc )

# Package "glmnetUtils" allows the use of R formulas for
# specifying glmnet models (as opposed to converting to matrices)
#library(devtools)
#install_github("hong-revo/glmnetUtils")
library(glmnetUtils)

# Package "broom" for tidy() function for pulling
# regression coefficients from glmnet models
library( broom )

# Package "rsample" for train/test split utilities
library( rsample )

```

```

Loading required package: Matrix
Loading required package: foreach
Loaded glmnet 2.0-18

```

```

Loading required package: lattice

```

```

Attaching package: mice

```

```

The following objects are masked from package:base:

```

```

  cbind, rbind

```

```

Registered S3 methods overwritten by 'ggplot2':

```

```

  method      from
  [.quosures   rlang
  c.quosures   rlang
  print.quosures rlang

```

```

Registered S3 method overwritten by 'rvest':

```

```

  method      from
  read_xml.response xml2

```

```

Attaching packages tidyverse 1.2.1

```

```

ggplot2 3.1.1      purrr   0.3.2
tibble  2.1.1      dplyr   0.8.1
tidyr   0.8.3      stringr 1.4.0
readr   1.3.1      forcats 0.4.0

```

```

Conflicts tidyverse_conflicts()

```

```

purrr::accumulate() masks foreach::accumulate()
tidyr::complete()   masks mice::complete()
tidyr::expand()     masks Matrix::expand()
dplyr::filter()     masks stats::filter()
dplyr::lag()        masks stats::lag()
purrr::when()       masks foreach::when()

```

Attaching package: skimr

The following object is masked from package:stats:

filter

Registered S3 method overwritten by 'GGally':

method from
+.gg ggplot2

Attaching package: GGally

The following object is masked from package:dplyr:

nasa

Attaching package: glmnetUtils

The following objects are masked from package:glmnet:

cv.glmnet, glmnet

2 Load data

- NYtowns dataset - marketing data
- [Product] Penetration is target variable
- Many other variables, mostly demographic data

```
In [2]: data <- read_excel( 'NYTown.xlsx' )
```

```
In [3]: data %>% dim
```

```
1.1006 2.248
```

```
In [4]: data <- data %>%  
  #select( -starts_with( 'Anc' ) ) %>%  
  select( -c( 'GEO_ID', "GEO_NAME", "NAME" ) )
```

3 Exploratory: univariate

- Distribution of Penetration is right skew and non-negative; not a count, but more like a rate

```
In [5]: # A Jupyter Notebook-specific directive  
  # sets the maximum number of rows to something above what is needed
```

```
# to show everything
options( repr.matrix.max.rows=300 )

In [6]: skim_to_wide( data ) %>% select( -type, -n )
```

variable <chr>	missing <chr>	complete <chr>	mean <chr>	sd <chr>	p0 <chr>	p25 <chr>
AncArab	0	1006	0.28	0.43	0	0
AncCzech	0	1006	0.49	0.5	0	0.1
AncDanish	0	1006	0.37	0.61	0	0.1
AncDutch	0	1006	4.16	2.87	0	2.3
AncEnglish	0	1006	14.37	5.39	0	10.6
AncFrCanad	0	1006	2.07	2.35	0	0.8
AncFrench	0	1006	6.58	6.47	0	3
AncGerman	0	1006	19.89	8.34	0	14.8
AncGreek	0	1006	0.34	0.49	0	0
AncHungary	0	1006	0.61	0.67	0	0.2
AncIrish	0	1006	17.73	5.12	0	14.3
AncItalian	0	1006	11.23	7.11	0	6.12
AncLithu	0	1006	0.26	0.34	0	0
AncNorweg	0	1006	0.67	0.62	0	0.2
AncOthr	0	1006	11.62	8.24	1.6	7.6
AncPolish	0	1006	5.91	4.57	0	3.1
AncPortug	0	1006	0.15	0.25	0	0
AncRussian	0	1006	0.93	1.26	0	0.2
AncScot	0	1006	2.37	1.32	0	1.5
AncScotIre	0	1006	1.39	0.8	0	0.9
AncSlovak	0	1006	0.27	0.57	0	0
AncSubSah	0	1006	0.12	0.26	0	0
AncSwedish	0	1006	1.35	2.35	0	0.5
AncSwiss	0	1006	0.43	0.64	0	0.1
AncUkraine	0	1006	0.68	0.83	0	0.2
AncUS	0	1006	6.85	3.66	0	4.1
AncWelsh	0	1006	1.22	1.42	0	0.5
AncWIndian	0	1006	0.25	0.91	0	0
AreaLand	0	1006	1.2e+08	9.6e+07	230812	8.2e
AreaWater	0	1006	9e+06	5e+07	0	1587
BadKitchen	0	1006	0.69	1.35	0	0.1
BadPlumbing	0	1006	0.75	1.35	0	0.2
BoatRVVan	0	1006	0.45	1.14	0	0
BornAfrica	14	992	1.26	3.36	0	0
BornAsia	14	992	17.28	17.1	0	0
BornEurope	14	992	49.57	23.92	0	32.6
BornLatAmer	14	992	14.48	18.74	0	0
BornNorAmer	14	992	16.79	20.57	0	2.1
BornOceania	14	992	0.62	2.82	0	0
BuiltAfter40	0	1006	67.87	13.23	28.8	59.3
BuiltAfter60	0	1006	51.62	13.99	10.7	43.4
BuiltAfter70	0	1006	40.81	13.1	5.3	33.0
BuiltAfter80	0	1006	26.13	9.98	1.3	19.7
BuiltAfter90	0	1006	13.22	5.93	0.2	9.2
BuiltAfter95	0	1006	6.43	3.62	0	4.1
BuiltAfter99	0	1006	1.48	1.21	0	0.6
CitizenNative	0	1006	96.41	4.61	53.9	95.9
CitizenNat	0	1006	1.97	2.12	0	0.7
CitizenNot	0	1006	1.62	2.75	0	0.3
CommuteAtHome	0	1006	4.45	2.96	0	2.6
CommuteAvgTravTime	0	1006	25.91	6.03	6	21.6

4 Optional: impute missing data

```
In [7]: data %>% is.na %>% sum
```

```
190
```

```
In [8]: tic()
```

```
  imputation_model <- mice( data, method='cart', m=1, maxit=1 )  
  toc()
```

```
iter imp variable
```

```
1    1  GPGuardP  BornEurope  BornAsia  BornAfrica  BornOceania  BornLatAmer  BornNorAmer  In
```

Warning message:

Number of logged events: 23

27.761 sec elapsed

```
In [9]: #glimpse( imputation_model )
```

```
In [10]: imputed_data <- mice::complete( imputation_model )
```

```
In [11]: imputed_data %>% is.na %>% sum
```

```
0
```

5 To evaluate predictions, split into train and test set

```
In [12]: set.seed( 42 )
```

```
  data_splitter <- initial_split( imputed_data, prop=0.8 )  
  train_data <- training( data_splitter )  
  test_data <- testing( data_splitter )
```

6 Perform LASSO

- Plain vanilla model (gaussian family w/identity link function)
 - Baseline: fitting linear model to poisson-style data
- Vary lambda regularization parameter, observe coefficient paths as lambda varies
 - By default, it will generate 100 different models with 100 different lambdas
 - A reasonable range of lambda vales are selected for you by default, also you can specify your own (or a range of lambdas)
- LASSO corresponds with glmnet with argument $\alpha = 1$ which is default
- glmnet standardized variables by default, so regression coefficients β_i is interpretable as relative importances

- Use dfmax argument to limit number of variables included in the model

```
In [13]: glmnet_lm_result <- glmnet( Penetration ~ ., data=train_data )
```

```
In [14]: print( glmnet_lm_result )
```

Call:

```
glmnet.formula(formula = Penetration ~ ., data = train_data)
```

Model fitting options:

Sparse model matrix: FALSE


Use model.frame: FALSE

Alpha: 1

Lambda summary:

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.000470	0.004705	0.047063	0.529144	0.470507	4.701200

```
In [15]: options(repr.plot.width=6, repr.plot.height=4)
plot( glmnet_lm_result, xvar='lambda', label=TRUE )
# can also use ggfortify version, which doesn't work very well with many variables
#options(repr.plot.width=30, repr.plot.height=30 )
#autoplot( glmnet_lm_result, label=FALSE )
#plot_glmnet( glmnet_lm_result, xvar='lambda', label=TRUE )
```



6.1 How many times does a given variable appear in a model?

```
In [16]: glmnet_lm_result %>%           # Take the glmnet result object
  coef %>%                             # get all the coefficients for all 100 models (l
  t %>%                                # transpose so the variables are in the columns
  as.matrix %>% as.data.frame %>%      # convert to something the tidyverse can manipulat
  map_int( ~ sum( . != 0 ) ) %>%      # For each column, return a count
  sort( decreasing = TRUE ) %>%      # sort from highest to lowest
  head( 40 ) %>%                      # show the top 20 used variables
  print
```

(Intercept)	HouseMultiFamily	AncItalian	ValueLT50K
100	98	97	95
BadPlumbing	NoCashRent	LessEq2Rooms	EduHSDip
94	94	92	91
JobAgriculture	MarFemaleDivorcees	WorkClassSelf	BuiltAfter90
89	88	87	86
AncSwiss	CommuteAtHome	PopRural	JobConstruct
85	82	82	81
BuiltAfter70	SchNurs	JobOfficeSales	EduColNoDeg
80	79	79	78
BadKitchen	CommuteAvgTravTime	CostDivIncLT15	VehicGT1
78	78	78	78
IndConstruction	Resid5Yrs	IndAgric	PovIndSeniors
77	76	75	75
IncAvgSocSec	IncRetirement	CostDivIncLT25	MarWidow
74	74	74	72
NativityOutUS	MortgageLT500	JobService	IncSupSec
72	72	71	71
AreaLand	SchKind	AncOthr	IndPublicAdmin
71	70	70	70

7 Which lambda gives the best model?

- BEST lambda is given by “lambda.min”: the at which the minimal MSE is achieved
- Use `cv.glmnet()` - Perform CROSS-VALIDATION LASSO
- Again, plain vanilla model (gaussian family w/identity link function)
- 10-fold cross validation to find optimal lambda parameter

```
In [17]: tic()
  glmnet_cv_result <- cv.glmnet( Penetration ~ ., data=train_data )
  toc()
```

1.736 sec elapsed

```
In [18]: print( glmnet_cv_result )
```


Call:

```
cv.glmnet.formula(formula = Penetration ~ ., data = train_data)
```

Model fitting options:

Sparse model matrix: FALSE

Use model.frame: FALSE

Number of crossvalidation folds: 10

Alpha: 1

Deviance-minimizing lambda: 0.2182104 (+1 SE): 0.6071833

```
In [19]: glmnet_cv_result$lambda.min
```


0.218210370672902

```
In [20]: options(repr.plot.width=6, repr.plot.height=5)
```

```
In [21]: log( glmnet_cv_result$lambda.min )
```

-1.52229567842249

```
In [22]: plot( glmnet_cv_result )
```



7.1 Given the best lambda, what are the regression coefficients?

- Remember, glmnet regression coefficients are standardized
 - Betas aren't in original units
 - Betas double as feature importances
 - To keep original units, use argument `standardize=False` in glmnet function call

```
In [23]: lm_coefs <- glmnet_cv_result %>%  
  coef( s = "lambda.min" ) %>%      # Get the betas from the best model  
  tidy %>%                          # Put betas into a data.frame  
  select( -column ) %>%             # prune unimportant column named "column"  
  arrange( desc( abs( value ) ) )   # Sort by absolute value
```

Warning message:

'tidy.dgCMatrix' is deprecated.

See help("Deprecated")Warning message:

'tidy.dgTMatrix' is deprecated.

See help("Deprecated")

```
In [24]: head( lm_coefs, 30 )
```

	row <chr>	value <dbl>
	(Intercept)	-18.37737222
	MarFemaleDivorcees	-5.11554208
	CostDivIncLT25	3.24460585
	CostDivIncLT15	2.86530431
	AncSwiss	0.98236846
	BadPlumbing	0.80539507
	LessEq2Rooms	0.54193077
	BadKitchen	0.32900049
	JobAgriculture	0.32006860
	WorkClassSelf	0.19349322
	IncRatioM2F	-0.12500977
	VehicGT1	0.09359990
	SchNurs	-0.08918499
A df[,2]: 30 CE 2	EduHSDip	0.08194796
	EduColNoDeg	-0.07702541
	IncSupSec	0.07538104
	NativityUS	0.07525924
	CommuteAvgTravTime	0.06972501
	IndConstruction	0.06312009
	CommuteAtHome	0.05860763
	BuiltAfter90	0.05768235
	NativityOutUS	-0.05674349
	IncRetirement	-0.05583241
	MarWidowedFemales	-0.05569634
	Resid5Yrs	0.05238722
	AncItalian	-0.05140025
	MarWidow	-0.05032510
	SchKind	-0.04929637
	MortgageLT700	0.04431094
	AncCzech	0.04281635

8 Use best lambda for prediction

- For functions like `predict()`, arg `s` is where you input lambda

```
In [25]: dim( train_data )
```

```
1. 805 2. 245
```

```
In [26]: lm_ypred <- predict(
  glmnet_cv_result, test_data, s="lambda.min" ) %>%
  as.numeric # convert predictions to R vector from R matrix
```

```
In [27]: length(lm_ypred)
```

```
201
```

```
In [28]: all_pred_results <- tibble( y=test_data$Penetration, glmnet_linear=lm_ypred )
```

8.1 Model achieves training R-squared of 0.72

```
In [29]: summary( lm( glmnet_linear ~ y, all_pred_results) )
```

Call:

```
lm(formula = glmnet_linear ~ y, data = all_pred_results)
```

Residuals:

Min	1Q	Median	3Q	Max
-17.4435	-2.0476	-0.1604	1.7866	17.6349

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	2.41597	0.39526	6.112	5.09e-09 ***
y	0.76420	0.03323	23.000	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 3.717 on 199 degrees of freedom


Multiple R-squared: 0.7266, Adjusted R-squared: 0.7253

F-statistic: 529 on 1 and 199 DF, p-value: < 2.2e-16

8.2 Make plot of predicted vs actual

- Note negative predictions!!!

```
In [30]: ggplot( all_pred_results, aes( y, glmnet_linear ) ) +  
         geom_point( alpha=0.3 )
```



8.3 Make log-log plot of predicted vs actual

- ggplot throws warning messages because it can't handle log transformations of negative predictions - it drops them from the figure.

```
In [31]: ggplot( all_pred_results, aes( y, glmnet_linear ) ) +
  geom_point( alpha=0.3 ) +
  scale_x_continuous(trans = 'log10') +
  scale_y_continuous(trans = 'log10') +
  annotation_logticks()
```


Warning message:

Transformation introduced infinite values in continuous x-axisWarning message in self\$trans\$tr

NaNs producedWarning message:

Transformation introduced infinite values in continuous y-axisWarning message:

Removed 12 rows containing missing values (geom_point).




9 Slightly better: Generalized Linear model

- Model the outcome variable Penetration as having a poisson distribution
- Use cross validation to find optimal lambda, then make predictions with that lambda.
- [GLM offsets](#) - advanced topic not covered here

```
In [32]: tic()
          glmnet_cv_result2 <- cv.glmnet( Penetration ~ ., data=train_data, family='poisson' )
          toc()
```

12.1 sec elapsed

```
In [33]: plot( glmnet_cv_result2 )
```



```
In [34]: lm2_ypred <- predict(
  glmnet_cv_result2, test_data, s="lambda.min", type = "response" ) %>%
  as.numeric # convert to R vector from R matrix
```

```
In [35]: length( lm2_ypred )
```

201

```
In [36]: all_pred_results <- all_pred_results %>%
  mutate( glmnet_poisson = lm2_ypred )
```

9.1 Poisson-style glmnet model R-squared is comparable to linear-style glmnet model

```
In [37]: summary( lm( glmnet_poisson ~ y, all_pred_results ) )
```

Call:

```
lm(formula = glmnet_poisson ~ y, data = all_pred_results)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-20.227	-1.217	-0.325	1.229	32.886

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1.63093	0.45776	3.563	0.000459 ***
y	0.86952	0.03848	22.596	< 2e-16 ***


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.305 on 199 degrees of freedom

Multiple R-squared: 0.7196, Adjusted R-squared: 0.7182

F-statistic: 510.6 on 1 and 199 DF, p-value: < 2.2e-16

```
In [38]: all_pred_results %>%  
  gather( key='model', value='Y_pred', -y ) %>%  
  ggplot( aes( y, Y_pred, color=model ) ) + geom_point()
```




```
In [39]: all_pred_results %>%
  gather( key='model', value='Y_pred', -y ) %>%
  ggplot( aes( y, Y_pred, color=model) ) + geom_point() +
  scale_x_continuous(trans = 'log10') +
  scale_y_continuous(trans = 'log10') +
  annotation_logticks()
```


Warning message:

Transformation introduced infinite values in continuous x-axisWarning message in self\$trans\$tr

NaNs producedWarning message:

Transformation introduced infinite values in continuous y-axisWarning message:

Removed 12 rows containing missing values (geom_point).



10 Use glmnet coefs for variable selection

```
In [40]: # remove row marked '(Intercept)'
interesting_linear_vars <- c( setdiff( lm_coefs[1:15,'row'], '(Intercept)'), 'Penetr
```

```
In [41]: interesting_linear_vars
```

```
1. 'MarFemaleDivorcees' 2. 'CostDivIncLT25' 3. 'CostDivIncLT15' 4. 'AncSwiss' 5. 'BadPlumbing' 6. 'LessEq2Rooms' 7. 'BadKitchen' 8. 'JobAgriculture' 9. 'WorkClassSelf' 10. 'IncRatioM2F' 11. 'VehicGT1' 12. 'SchNurs' 13. 'EduHSDip' 14. 'EduColNoDeg' 15. 'Penetration'
```

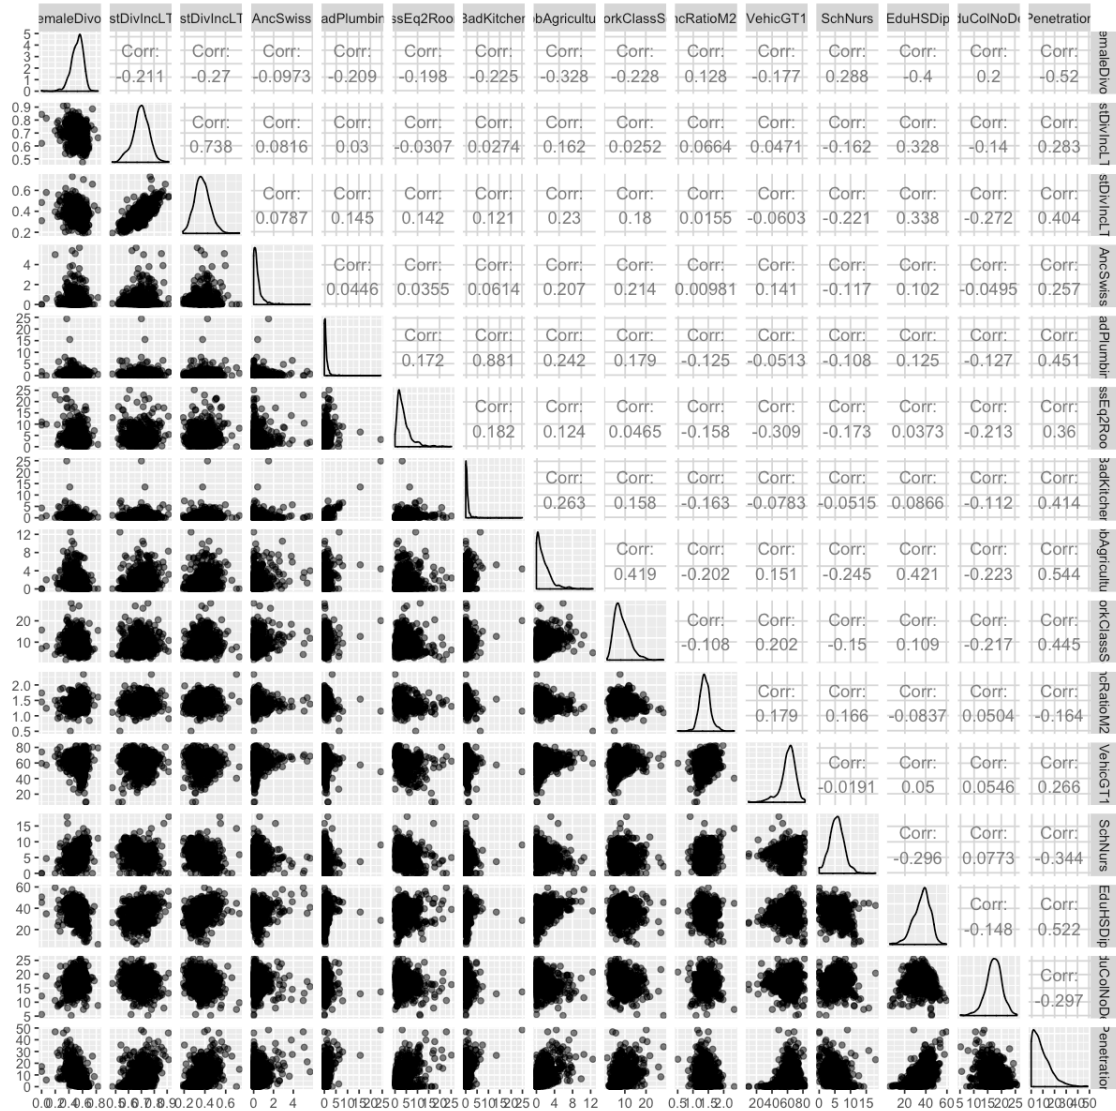
```
In [42]: train_data_subset <- train_data %>% select( interesting_linear_vars )
```

10.1 Exploratory: Univariate and Bivariate dists

```
In [43]: skim_to_wide( train_data_subset ) %>% select( -type, -n )
```

	variable <chr>	missing <chr>	complete <chr>	mean <chr>	sd <chr>	p0 <chr>	p25 <chr>	p50 <chr>	p75 <chr>
A tibble: 15 × 11	AncSwiss	0	805	0.42	0.62	0	0.1	0.2	0.5
	BadKitchen	0	805	0.68	1.3	0	0.1	0.4	0.9
	BadPlumbing	0	805	0.74	1.33	0	0.2	0.4	0.9
	CostDivIncLT15	0	805	0.37	0.075	0.19	0.33	0.37	0.4
	CostDivIncLT25	0	805	0.7	0.063	0.47	0.66	0.7	0.7
	EduColNoDeg	0	805	17.24	2.92	5.2	15.6	17.4	19
	EduHSDip	0	805	36.23	7.89	6	31.4	37.2	41
	IncRatioM2F	0	805	1.39	0.18	0.5	1.28	1.38	1.5
	JobAgriculture	0	805	1.54	1.66	0	0.3	1.1	2.2
	LessEq2Rooms	0	805	3.84	3.34	0	1.7	3	4.8
	MarFemaleDivorcees	0	805	0.52	0.089	0	0.47	0.53	0.5
	Penetration	0	805	8.66	7.9	0	2.5	6.7	13
	SchNurs	0	805	5.58	2.36	0	4.1	5.6	7
	VehicGT1	0	805	60.93	9.87	9.6	56.8	62.3	67
	WorkClassSelf	0	805	8.61	3.65	1.6	6	7.9	10

```
In [44]: options(repr.plot.width=10, repr.plot.height=10)
          ggpairs( train_data_subset, aes( alpha=0.001 ) )
```



11 Regular Linear Model

```
In [45]: model0 <- lm( Penetration ~ ., data=train_data_subset )
```

```
In [46]: summary( model0 )
```

Call:

```
lm(formula = Penetration ~ ., data = train_data_subset)
```

Residuals:

Min	1Q	Median	3Q	Max
-12.3993	-2.3369	-0.2804	1.9329	15.2693

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-22.52440	2.93277	-7.680	4.70e-14	***
MarFemaleDivorcees	-7.50162	1.89528	-3.958	8.24e-05	***
CostDivIncLT25	5.32176	3.48363	1.528	0.127000	
CostDivIncLT15	11.31104	3.06880	3.686	0.000244	***
AncSwiss	0.98030	0.23534	4.165	3.45e-05	***
BadPlumbing	1.26269	0.22902	5.514	4.76e-08	***
LessEq2Rooms	0.74801	0.04796	15.597	< 2e-16	***
BadKitchen	0.38757	0.23727	1.633	0.102774	
JobAgriculture	0.41190	0.10843	3.799	0.000157	***
WorkClassSelf	0.37321	0.04526	8.246	6.81e-16	***
IncRatioM2F	-2.49474	0.85991	-2.901	0.003821	**
VehicGT1	0.24610	0.01651	14.906	< 2e-16	***
SchNurs	-0.15866	0.06578	-2.412	0.016098	*
EduHSDip	0.29024	0.02195	13.223	< 2e-16	***
EduColNoDeg	-0.13671	0.05213	-2.623	0.008895	**

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

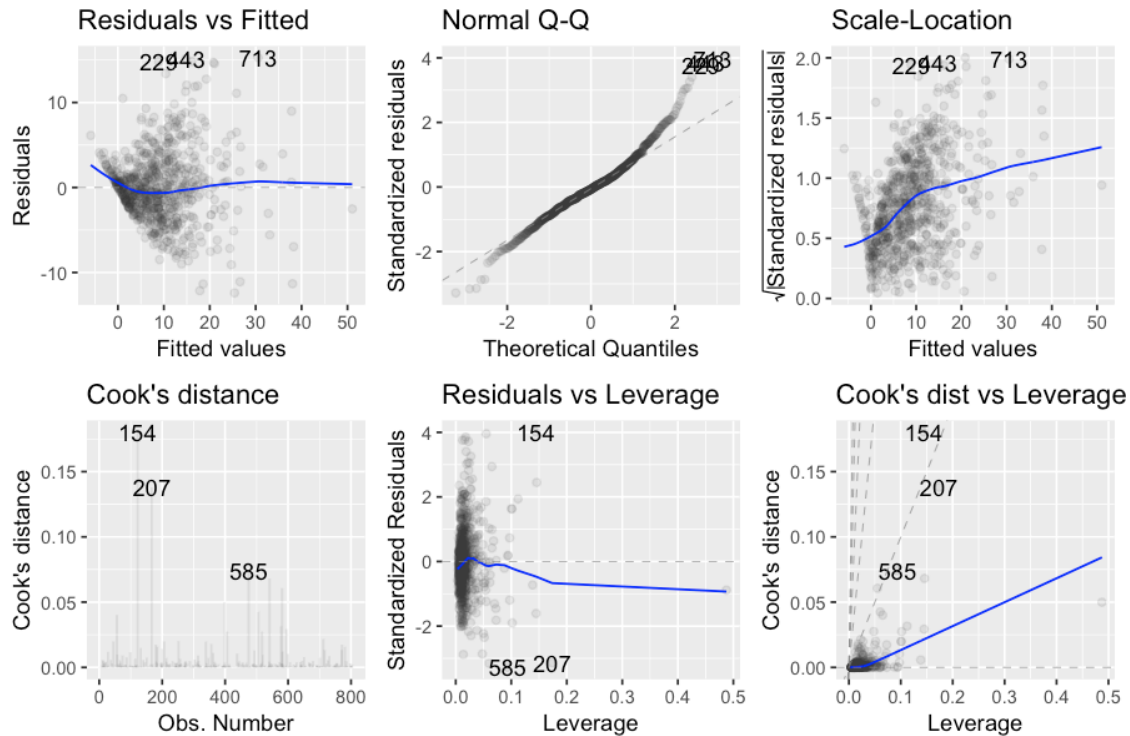
Residual standard error: 3.969 on 790 degrees of freedom

Multiple R-squared: 0.7519, Adjusted R-squared: 0.7475

F-statistic: 171 on 14 and 790 DF, p-value: < 2.2e-16

```
In [47]: options( repr.plot.width=7.5, repr.plot.height=5 )
```

```
In [48]: autoplot( model0, which = 1:6, ncol=3, alpha=0.1)
```




```
In [49]: all_pred_results <- all_pred_results %>%
         mutate( lm = predict( model0, test_data ) )
```

```
In [50]: all_pred_results %>%
         gather( key='model', value='Y_pred', -y ) %>%
         ggplot( aes( y, Y_pred, color=model ) ) + geom_point() #+
         #scale_x_continuous(trans = 'log10') +
         #scale_y_continuous(trans = 'log10') +
         #annotation_logticks()
```

Warning message:

attributes are not identical across measure variables;
they will be dropped



12 GLM model

```
In [51]: glm_coefs <- glmnet_cv_result2 %>%
  coef( s = "lambda.min" ) %>%      # Get the betas from the best model
  tidy %>%                          # Put betas into a data.frame
  select( -column ) %>%             # prune unimportant column named "column"
  arrange( desc( abs( value ) ) )   # Sort by absolute value
```

Warning message:

'tidy.dgCMatrix' is deprecated.

See help("Deprecated")Warning message:

'tidy.dgTMatrix' is deprecated.

See help("Deprecated")

```
In [52]: glm_coefs %>% head(15)
```

	row <chr>	value <dbl>
A df[,2]: 15 CE 2	MarFemaleDivorcees	-0.38459635
	CostDivIncLT25	0.38011921
	LabFCivilEmployed	0.25292790
	(Intercept)	0.21780932
	AncSubSah	-0.13735906
	CostDivIncLT15	0.13313636
	AncSwiss	0.04909032
	BadKitchen	0.02717750
	AncItalian	-0.02633768
	HouseMultiFamily	-0.02609654
	AncCzech	0.02494139
	BadPlumbing	0.02413368
	WorkClassSelf	0.02033181
	JobAgriculture	0.01565186
	IncSupSec	0.01512092

```
In [53]: # remove row marked '(Intercept)'
         interesting_poisson_vars <- c( setdiff( glm_coefs[1:15,'row'], '(Intercept)'), 'Penetration' )
```

```
In [54]: data_subset <- train_data %>% select( interesting_poisson_vars )
```

```
In [55]: model1 <- glm( Penetration ~ ., data=data_subset, family='quasipoisson' )
```

```
In [56]: summary( model1 )
```

Call:

```
glm(formula = Penetration ~ ., family = "quasipoisson", data = data_subset)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-5.0784	-1.0049	-0.2477	0.6750	4.6136

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	2.499317	0.457565	5.462	6.30e-08	***
MarFemaleDivorcees	-1.028500	0.178225	-5.771	1.13e-08	***
CostDivIncLT25	0.848191	0.362827	2.338	0.019650	*
LabFCivilEmployed	-0.207178	0.420219	-0.493	0.622133	
AncSubSah	-0.357826	0.117860	-3.036	0.002476	**
CostDivIncLT15	-0.004742	0.315302	-0.015	0.988005	
AncSwiss	0.085180	0.018595	4.581	5.38e-06	***
BadKitchen	-0.009841	0.019379	-0.508	0.611718	
AncItalian	-0.043466	0.003932	-11.053	< 2e-16	***
HouseMultiFamily	-0.044923	0.002607	-17.233	< 2e-16	***
AncCzech	0.099778	0.029782	3.350	0.000846	***
BadPlumbing	0.045596	0.019257	2.368	0.018138	*

```

WorkClassSelf      0.023411    0.004522    5.177 2.86e-07 ***
JobAgriculture     0.053187    0.009138    5.820 8.53e-09 ***
IncSupSec          0.041014    0.007111    5.768 1.15e-08 ***

```

Signif. codes: 0 *** 0.001 ** 0.01 * 0.05 . 0.1 1

(Dispersion parameter for quasipoisson family taken to be 1.740372)

```

Null deviance: 5592.2  on 804  degrees of freedom
Residual deviance: 1424.6  on 790  degrees of freedom
AIC: NA

```

Number of Fisher Scoring iterations: 5

```
In [57]: glm_preds <- as.numeric( predict( model1, test_data, type='response' ) )
```

```
In [58]: all_pred_results <- all_pred_results %>%
         mutate( glm = glm_preds )
```

```
In [59]: summary( lm( glm ~ y, all_pred_results ) )
```

Call:

```
lm(formula = glm ~ y, data = all_pred_results)
```

Residuals:

```

      Min       1Q   Median       3Q      Max
-13.8162  -2.4198  -0.9994   1.2441  22.4002

```

Coefficients:

```

              Estimate Std. Error t value Pr(>|t|)
(Intercept)  3.10600     0.49328   6.297 1.91e-09 ***
y            0.69032     0.04147  16.648 < 2e-16 ***

```

Signif. codes: 0 *** 0.001 ** 0.01 * 0.05 . 0.1 1


Residual standard error: 4.639 on 199 degrees of freedom

Multiple R-squared: 0.5821, Adjusted R-squared: 0.58

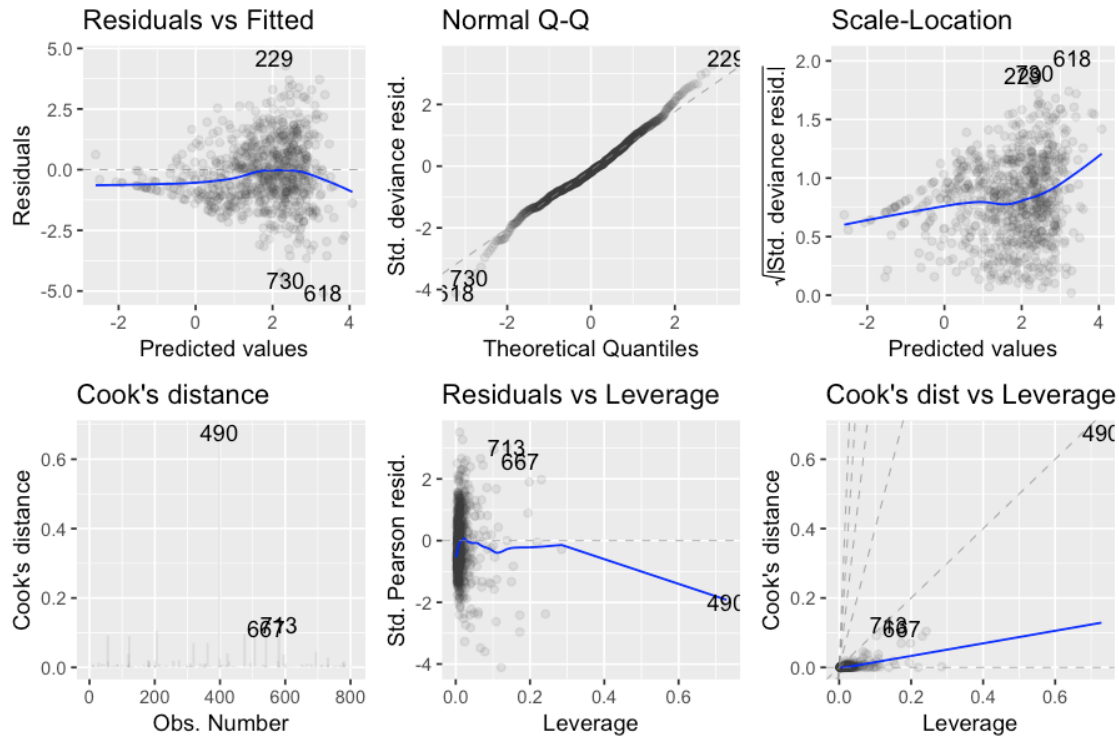
F-statistic: 277.1 on 1 and 199 DF, p-value: < 2.2e-16

```
In [60]: all_pred_results %>%
         gather( key='model', value='Y_pred', -y ) %>%
         ggplot( aes( y, Y_pred, color=model ) ) + geom_point() #+
         #scale_x_continuous(trans = 'log10') +
         #scale_y_continuous(trans = 'log10') +
         #annotation_logticks()
```


Warning message:
attributes are not identical across measure variables;
they will be dropped



```
In [61]: autoplot( model1, which = 1:6, ncol=3, alpha=0.1)
```



12.1 What about log transforming the skewed vars?

```
In [ ]: log_transform_vars <- c( "MortgageLT500", "CommuteAtHome", 'BadPlumbing', "BoatRVVan",

In [ ]: transformed_data_subset <- data_subset %>%
      mutate_at( log_transform_vars, ~ log( . + 1) )

In [ ]: skim_to_wide( transformed_data_subset )

In [ ]: skim_to_wide( data_subset )

In [ ]: model2 <- glm( Penetration ~ ., data=transformed_data_subset, family='quasipoisson' )

In [ ]: summary( model2 )

In [ ]: autoplot( model2, which = 1:6, ncol=3)

In [ ]: anova( model1, model2)
```